

Natural Language Processing


Sequence to Sequence Models and Attention

Felipe Bravo-Marquez

June 6, 2023

Language Models and Language Generation

- Language modeling is the task of assigning a probability to sentences in a language.
- Example: what is the probability of seeing the sentence “the lazy dog barked loudly”?
- The task can be formulated as the task of predicting the probability of seeing a word conditioned on previous words:

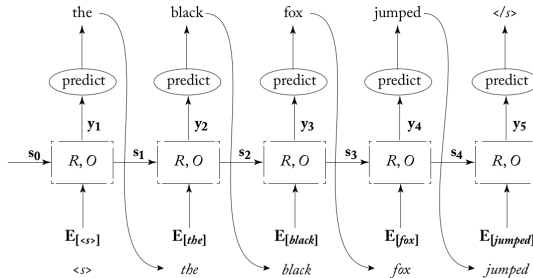

$$P(w_i | w_1, w_2, \dots, w_{i-1}) = \frac{P(w_1, w_2, \dots, w_{i-1}, w_i)}{P(w_1, w_2, \dots, w_{i-1})}$$

Language Models and Language Generation

- RNNs can be used to train language models by tying the output at time i with its input at time $i + 1$.
- This network can be used to generate sequences of words or random sentences.
- Generation process: predict a probability distribution over the first word conditioned on the start symbol, and draw a random word according to the predicted distribution.
- Then predict a probability distribution over the second word conditioned on the first, and so on, until predicting the end-of-sequence $< /s >$ symbol.

Language Models and Language Generation

- After predicting a distribution over the next output symbols $P(t_i = k | t_{1:i-1})$, a token t_i is chosen and its corresponding embedding vector is fed as the input to the next step.



- Teacher-forcing: during **training** the generator is fed with the ground-truth previous word even if its own prediction put a small probability mass on it.
- It is likely that the generator would have generated a different word at this state in **test time**.

Sequence to Sequence Problems

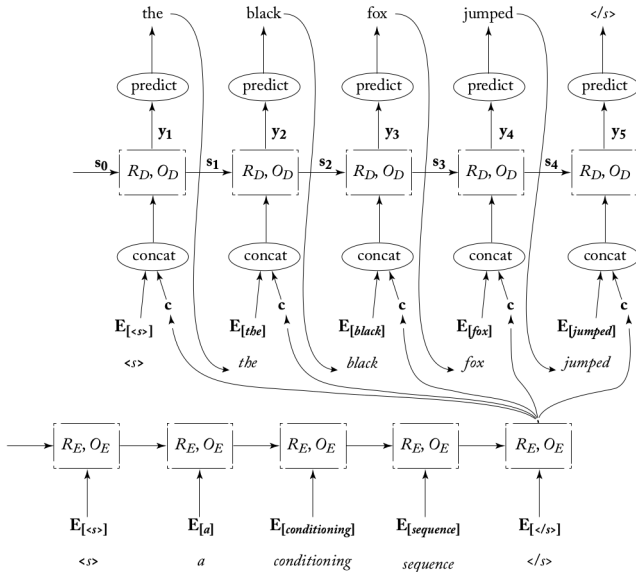
Nearly any task in NLP can be formulated as a sequence to sequence (or conditioned generation) task i.e., **generate output sequences from input ones**. Input and output sequences can have different lengths.

- **Machine Translation:** source language to target language.
- **Summarization:** long text to short text.
- **Dialogue (chatbots):** previous utterances to next utterance.


Conditioned Generation

- While using the RNN as a generator is a cute exercise for demonstrating its strength, the power of RNN generator is really revealed when moving to a conditioned generation or **encoder-decoder framework**.
- Core idea: **using two RNNs.**
- **Encoder:** One RNN is used to **encode** the source **input** into a vector \vec{c} .
- **Decoder:** Another RNN is used to **decode the encoder's output and generate the target output.**
- At each stage of the generation process the context vector \vec{c} is concatenated to the input \hat{t}_j and the concatenation is fed into the RNN.

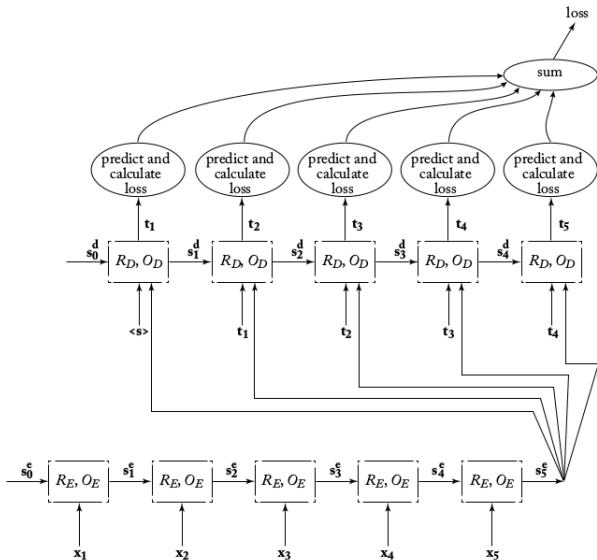
Encoder Decoder Framework




Conditioned Generation

- This setup is useful for mapping sequences of length n to sequences of length m .
 - The encoder summarizes the source sentence as a vector \vec{c} .
 - The decoder RNN is then used to predict (using a language modeling objective) the target sequence words conditioned on the previously predicted words as well as the encoded sentence \vec{c} .
 - The encoder and decoder RNNs are trained jointly.
 - The supervision happens only for the decoder RNN, but the gradients are propagated all the way back to the encoder RNN.
- 

Sequence to Sequence Training Graph




Neural Machine Translation



Kyunghyun Cho et al. 2014

Machine Translation BLEU progress over time



[Edinburgh En-De WMT]

⁰source: http://www.meta-net.eu/events/meta-forum-2016/slides/09_sennrich.pdf

Decoding Approaches

- The decoder aims to generate the output sequence with maximal score (or maximal probability), i.e., such that $\sum_{i=1}^n P(\hat{t}_i | \hat{t}_{1:i-1})$ is maximized.
- The non-markovian nature of the RNN means that the probability function cannot be decomposed into factors that allow for exact search using standard dynamic programming.
- Exact search: finding the optimum sequence requires evaluating every possible sequence (computationally prohibitive).
- Thus, it only makes sense to solving the optimization problem above approximately.
- Greedy search: choose the highest scoring prediction (word) at each step.
- This may result in sub-optimal overall probability leading to prefixes that are followed by low-probability events.

Greedy Search




Figure 6.4: (a) Search space depicted as a tree. (b) Greedy search.

Beam Search

- Beam search interpolates between the exact search and the greedy search by changing the size K of hypotheses maintained throughout the search procedure [Cho, 2015].
- The Beam search algorithm works in stages.
- We first pick the K starting words with the highest probability
- At each step, each candidate sequence is expanded with all possible next steps.
- Each candidate step is scored.
- The K sequences with the most likely probabilities are retained and all other candidates are pruned.
- The search process can halt for each candidate separately either by reaching a maximum length, by reaching an end-of-sequence token, or by reaching a threshold likelihood.
- The sentence with the highest overall probability is selected.

Conditioned Generation with Attention

- In the encoder-decoder networks the input sentence is encoded into a single vector, which is then used as a conditioning context for an RNN-generator.
- This architecture forces the encoded vector \vec{c} to contain all the information required for generation.
- It doesn't work well for long sentences!
- It also requires the generator to be able to extract this information from the fixed-length vector.
- “You can't cram the meaning of a whole sentence into a single vector!” -Raymond Mooney
- This architecture can be substantially improved (in many cases) by the addition of an attention mechanism.
- The attention mechanism attempts to solve this problem by allowing the decoder to “look back” at the encoder's hidden states based on its current state.

Conditioned Generation with Attention

- The input sentence (a length n input sequence $\vec{x}_{1:n}$) is encoded using a biRNN as a sequence of vectors $\vec{c}_{1:n}$.
- The decoder uses a soft attention mechanism in order to decide on which parts of the encoding input it should focus.
- At each stage j the decoder sees a weighted average of the vectors $\vec{c}_{1:n}$, where the attention weights ($\vec{\alpha}^j$) are chosen by the attention mechanism.

$$\vec{c}^j = \sum_{i=1}^n \alpha_{[i]}^j \cdot \vec{c}_i$$

- The elements of $\vec{\alpha}^j$ are all positive and sum to one.

Conditioned Generation with Attention

- Unnormalized attention weights ($\tilde{\alpha}_{[j]}^j$) are produced taking into account the decoder state at time j (\vec{s}_j) and each of the vectors \vec{c}_i .
- They can be obtained in various ways, basically any differentiable function returning a scalar out of two vectors \vec{s}_j and \vec{c}_i could be employed.
- The simplest approach is a dot product: $\tilde{\alpha}_{[j]}^j = \vec{s}_j \cdot \vec{c}_i$.
- The one we will use in these slides is Additive attention, which uses a Multilayer Perceptron: $\tilde{\alpha}_{[j]}^j = MLP^{att}([\vec{s}_j; \vec{c}_i]) = \vec{v} \cdot \tanh([\vec{s}_j; \vec{c}_i]U + \vec{b})$

Conditioned Generation with Attention

- These unnormalized weights are then normalized into a probability distribution using the softmax function.

$$\text{attend}(c_{1:n}, \hat{t}_{1:j}) = c^j$$


$$c^j = \sum_{i=1}^n \alpha_{[i]}^j \cdot c_i$$

$$\alpha^j = \text{softmax}(\bar{\alpha}_{[1]}^j, \dots, \bar{\alpha}_{[n]}^j)$$

$$\bar{\alpha}_{[i]}^j = \text{MLP}^{\text{att}}([s_j; c_i]),$$

- The encoder, decoder, and attention mechanism are all trained jointly in order to play well with each other.

Attention



Conditioned Generation with Attention

The entire sequence-to-sequence generation with attention is given by:

$$p(t_{j+1} = k \mid \hat{t}_{1:j}, \mathbf{x}_{1:n}) = f(O_{\text{dec}}(s_{j+1}))$$

$$s_{j+1} = R_{\text{dec}}(s_j, [\hat{t}_j; c^j])$$

$$c^j = \sum_{i=1}^n \alpha_{[i]}^j \cdot c_i$$

$$c_{1:n} = \text{biRNN}_{\text{enc}}^*(\mathbf{x}_{1:n})$$

$$\alpha^j = \text{softmax}(\bar{\alpha}_{[1]}^j, \dots, \bar{\alpha}_{[n]}^j)$$

$$\bar{\alpha}_{[i]}^j = \text{MLP}^{\text{att}}([s_j; c_i])$$

$$\hat{t}_j \sim p(t_j \mid \hat{t}_{1:j-1}, \mathbf{x}_{1:n})$$

$$f(z) = \text{softmax}(\text{MLP}^{\text{out}}(z))$$

$$\text{MLP}^{\text{att}}([s_j; c_i]) = \mathbf{v} \tanh([s_j; c_i]U + \mathbf{b}).$$

Conditioned Generation with Attention

- Why use the biRNN encoder to translate the conditioning sequence $\vec{x}_{1:n}$ into the context vectors $\vec{c}_{1:n}$?
- Why we just don't attend directly on the inputs (word embeddings)
 $MLP^{att}([\vec{s}_j; \vec{x}_i])$?
- We could, but we get important **benefits from the encoding process.**
- First, the biRNN vectors \vec{c}_i represent the items \vec{x}_i in their sentential context.
- **Sentential context: a window focused around the input item \vec{x}_i and not the item itself.**
- **Second, by having a trainable encoding component that is trained jointly with the decoder, the encoder and decoder evolve together.**
- Hence, the network can learn to encode relevant properties of the input that are useful for decoding, and that may not be present at the source sequence $\vec{x}_{1:n}$ directly.

Attention and Word Alignments

- In the context of machine translation, one can think of MLP^{att} as computing a soft alignment between the current decoder state \vec{s}_j (capturing the recently produced foreign words) and each of the source sentence components \vec{c}_i .




Fig. 2. Visualization of the attention weights α_j^t of the attention-based neural machine translation model [32]. Each row corresponds to the output symbol, and each column the input symbol. Brighter the higher α_j^t .

Figure: Source: [Cho et al., 2015]

Other types of Attention

Summary

Below is a summary table of several popular attention mechanisms (or broader categories of attention mechanisms).

Name	Alignment score function	Citation
Additive(*)	$\text{score}(s_t, h_i) = \mathbf{v}_a^\top \tanh(\mathbf{W}_a[s_t; h_i])$	Bahdanau2015
Location-Based	$\alpha_{t,i} = \text{softmax}(\mathbf{W}_a s_t)$ Note: This simplifies the softmax alignment max to only depend on the target position.	Luong2015
General	$\text{score}(s_t, h_i) = s_t^\top \mathbf{W}_a h_i$ where \mathbf{W}_a is a trainable weight matrix in the attention layer.	Luong2015
Dot-Product	$\text{score}(s_t, h_i) = s_t^\top h_i$	Luong2015
Scaled Dot-Product(^)	$\text{score}(s_t, h_i) = \frac{s_t^\top h_i}{\sqrt{n}}$ Note: very similar to the dot-product attention except for a scaling factor; where n is the dimension of the source hidden state.	Vaswani2017
Self-Attention(&)	Relating different positions of the same input sequence. Theoretically the self-attention can adopt any score functions above, but just replace the target sequence with the same input sequence.	Cheng2016
Global/Soft	Attending to the entire input state space.	Xu2015
Local/Hard	Attending to the part of input state space; i.e. a patch of the input image.	Xu2015 ; Luong2015

(*) Referred to as "concat" in Luong, et al., 2015 and as "additive attention" in Vaswani, et al., 2017.

(^) It adds a scaling factor $1/\sqrt{n}$, motivated by the concern when the input is large, the softmax function may have an extremely small gradient, hard for efficient learning.

(&) Also, referred to as "intra-attention" in Cheng et al., 2016 and some other papers.

Figure: Source: <https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>

Questions?

Thanks for your Attention!

References I



Cho, K. (2015).

Natural language understanding with distributed representation.

arXiv preprint arXiv:1511.07916.



Cho, K., Courville, A., and Bengio, Y. (2015).

Describing multimedia content using attention-based encoder-decoder networks.

IEEE Transactions on Multimedia, 17(11):1875–1886.